

CT5102: Programming for Data Analytics

Week 3: Matrices and Data Frames

<https://github.com/JimDuggan/CT5102>

Dr. Jim Duggan,
Information Technology,
School of Engineering & Informatics

One-slide summary to date...

- **Vectors** – fundamental data type in R. *Vectorization*. “No loops required!”
- **Functions** – modularity, divide and conquer, reuse. *Apply family*
- **Lists** – complex data structures

Input Vector

1
4
9
16
25

sqrt()

Output Vector

1
2
3
4
5

```
test<-function(n1,n2=20){  
  n1+n2  
}
```

Matrices

- A matrix is a vector that has a number of rows and a number of columns, as well as having defined modes (Matloff 2009).
- Matrix subscripts, similar to vectors, commence at 1, and there are useful operations to access row and column elements.
- A matrix can be initialized from a vector, where the numbers of rows and columns are specified.
- R stores matrices by column-major order, and by default matrices are filled in this manner.

Examples

```
v<-c(10,20,30,40,50,60)

m1<-matrix(v,nrow=3,ncol=2)

m2<-matrix(v,nrow=3,ncol=2,
            byrow=T)
```

```
> v
[1] 10 20 30 40 50 60
> m1
      [,1] [,2]
[1,]   10   40
[2,]   20   50
[3,]   30   60
> m2
      [,1] [,2]
[1,]   10   20
[2,]   30   40
[3,]   50   60
```

Accessing Elements

```
> m1
      [,1] [,2]
[1,]   10  40
[2,]   20  50
[3,]   30  60
> m1[1,1]
[1] 10
> m1[1,]
[1] 10 40
> m1[1:2,]
      [,1] [,2]
[1,]   10  40
[2,]   20  50
```

```
> m1
      [,1] [,2]
[1,]   10  40
[2,]   20  50
[3,]   30  60
> m1[,1]
[1] 10 20 30
> m1[,2]
[1] 40 50 60
> m1[,-(1)]
[1] 40 50 60
```

Useful Functions for Matrices

```
> m1
      [,1] [,2]
[1,]   10   40
[2,]   20   50
[3,]   30   60
> dim(m1)
[1] 3 2
>
> nrow(m1)
[1] 3
> ncol(m1)
[1] 2
```

```
> m1
      [,1] [,2]
[1,]   10   40
[2,]   20   50
[3,]   30   60
> rowSums(m1)
[1] 50 70 90
> colSums(m1)
[1] 60 150
>
> rowMeans(m1)
[1] 25 35 45
> colMeans(m1)
[1] 20 50
```

Filtering Matrices

- Filtering can also be performed on matrices.
- For example, if a query is required to find all rows that have *column 1 values greater than 20*.
- First a logical vector could be applied to the full column with the specified condition

```
> m1
      [,1] [,2]
[1,]   10  40
[2,]   20  50
[3,]   30  60
> b1<-m1[,1] > 20
> b1
[1] FALSE FALSE  TRUE
> m1[b1,]
[1] 30 60
  1
```

Appending rows and columns

```
> m1
      [,1] [,2]
[1,]   10   40
[2,]   20   50
[3,]   30   60
> rbind(m1,c(40,70))
      [,1] [,2]
[1,]   10   40
[2,]   20   50
[3,]   30   60
[4,]   40   70
```

```
> m1
      [,1] [,2]
[1,]   10   40
[2,]   20   50
[3,]   30   60
> cbind(m,c(70,80,90))
      [,1] [,2] [,3]
[1,]   10   40   70
[2,]   20   50   80
[3,]   30   60   90
```


Sample Matrix Operations

Operator or Function	Description
$A * B$	Element-wise multiplication
A / B	Element-wise division
$A \%*\% B$	Matrix multiplication
$t(A)$	Transpose of A
$e<-eigen(A)$	List of eigenvalues and eigenvectors for matrix A

apply() function

- The **apply()** function can be used to process rows and columns for a matrix, and the general form of this function (Matloff 2009) is **apply(m, dimcode, f, fargs)**, where:
 - **m** is the target matrix
 - **dimcode** identifies whether it's a row or column target. The number 1 applies to rows, whereas 2 applies to columns
 - **f** is the function to be called
 - **fargs** are the optional set of arguments that can be applied to the function **f**.

apply() in action

```
> m1
      [,1] [,2]
[1,]   10   40
[2,]   20   50
[3,]   30   60
> # 1 applies to rows
> apply(m1,1,min)
[1] 10 20 30
```

```
> m1
      [,1] [,2]
[1,]   10   40
[2,]   20   50
[3,]   30   60
> # 2 applies to columns
> apply(m1,2,min)
[1] 10 40
```

Row and Column Names

```
> dimnames(m1) <- list(rownames(m1, do.NULL = FALSE, prefix = "row"),  
+                       colnames(m1, do.NULL = FALSE, prefix = "col"))  
> m1
```

	col1	col2
row1	10	40
row2	20	50
row3	30	60

Challenge 3.1

- Create a 5x5 square matrix of random uniform numbers
- Add a new column that contains the maximum value in each row
- Add a new row that contains the maximum value in each column

Challenge 3.2

Customer Preference Matrix

- Write a function that creates an $n \times m$ matrix of customers and products, and for each customer, assigns a random preference vector showing their preference for a product.
- Customers should express a preference value for a random $x\%$ of products
- The function default values should be 10, 10 and 20%
- Preferences should be generated on a uniform scale from 1 to 5 (representing 1 to 5 stars)
- Columns should be named Pr1,Pr2,...,PrN
- Rows should be names Cust1, Cust2, ..., CustN

Data Frames

- A data frame is similar to a matrix, as it has a two-dimensional rows and columns structure
- Differs from a matrix in that each column can have a different mode (Matloff 2009).
- Can be viewed as a set of vectors, organised into a column format

```
> ids<-c("1234567","1234568")
> fName<-c("Jane","Matt")
> sName<-c("Smith","Johnson")
> ages<-c(21,25)
> ids
[1] "1234567" "1234568"
> fName
[1] "Jane" "Matt"
> sName
[1] "Smith" "Johnson"
> ages
[1] 21 25
```

Creating a Data Frame

```
s<-data.frame(ID=ids,FirstName=fNames,Surname=sNames,  
              Age=ages,stringsAsFactors=FALSE)
```

```
> s
```

	ID	FirstName	Surname	Age
1	1234567	Jane	Smith	21
2	1234568	Matt	Johnson	25

```
> str(s)
```

```
'data.frame':  2 obs. of  4 variables:  
 $ ID          : chr  "1234567" "1234568"  
 $ FirstName: chr  "Jane" "Matt"  
 $ Surname   : chr  "Smith" "Johnson"  
 $ Age       : num  21 25
```


Accessing Row Data (matrix method)

```
> s
```

	ID	FirstName	Surname	Age
1	1234567	Jane	Smith	21
2	1234568	Matt	Johnson	25

```
> s[1,]
```

	ID	FirstName	Surname	Age
1	1234567	Jane	Smith	21

```
> s[1:2,]
```

	ID	FirstName	Surname	Age
1	1234567	Jane	Smith	21
2	1234568	Matt	Johnson	25

```
> s[-1,]
```

	ID	FirstName	Surname	Age
2	1234568	Matt	Johnson	25

Accessing Column Data (matrix method)

```
> s
      ID FirstName Surname Age
1 1234567      Jane   Smith  21
2 1234568      Matt Johnson  25
> s[, "Age"]
[1] 21 25
> s[, 1]
[1] "1234567" "1234568"
> s[, 2]
[1] "Jane" "Matt"
> s[, 3]
[1] "Smith"  "Johnson"
> s[, 4]
[1] 21 25
```

Accessing Column Data (using tags)

```
> s
      ID FirstName Surname Age
1 1234567      Jane   Smith  21
2 1234568      Matt Johnson  25
> s$ID
[1] "1234567" "1234568"
> s$FirstName
[1] "Jane" "Matt"
> s$Age[1:2]
[1] 21 25
> max(s$Age)
[1] 25
|
```

Subsetting Row Data (Matrix method)

```
> s
```

	ID	FirstName	Surname	Age
1	1234567	Jane	Smith	21
2	1234568	Matt	Johnson	25

```
> sub<-s[s$Age>21,]
```

```
> sub
```

	ID	FirstName	Surname	Age
2	1234568	Matt	Johnson	25

|

subset() function

```
> s
```

	ID	FirstName	Surname	Age
1	1234567	Jane	Smith	21
2	1234568	Matt	Johnson	25

```
> sub<-subset(s,s$Age>21)
```

```
> sub
```

	ID	FirstName	Surname	Age
2	1234568	Matt	Johnson	25

Adding extra information to a data frame

```
> s
```

	ID	FirstName	Surname	Age
1	1234567	Jane	Smith	21
2	1234568	Matt	Johnson	25

```
> s$Discount<-ifelse(s$Age<=21,0.25,0.10)
```

```
> s
```

	ID	FirstName	Surname	Age	Discount
1	1234567	Jane	Smith	21	0.25
2	1234568	Matt	Johnson	25	0.10

Merging Data Frames

- For data analytics, opportunities often arise by merging different data sets into a single data frame, and the **merge()** function facilitates this
- In our student example, we could have a second data frame that stores examination results.

```
> ids<-c("1234567","1234568")
> subjects<-c("CT111","CT111")
> grade<-c(80,80)
> res<-data.frame(ID=ids,Subject=subjects,Grade=grade,stringsAsFactors=FALSE)
> res
```

	ID	Subject	Grade
1	1234567	CT111	80
2	1234568	CT111	80

1

Merging Code

```
> s
```

	ID	FirstName	Surname	Age	Discount
1	1234567	Jane	Smith	21	0.25
2	1234568	Matt	Johnson	25	0.10

```
> res
```

	ID	Subject	Grade
1	1234567	CT111	80
2	1234568	CT111	80


```
> new<-merge(s,res,by="ID")
```

```
> new
```

	ID	FirstName	Surname	Age	Discount	Subject	Grade
1	1234567	Jane	Smith	21	0.25	CT111	80
2	1234568	Matt	Johnson	25	0.10	CT111	80

Accessing Real Data Sets

Year	Leinster	Munster	Connacht	Ulster
1841	1973731	2396161	1418859	740048
1851	1672738	1857736	1010031	571052
1861	1457635	1513558	913135	517783
1871	1339451	1393485	846213	474038
1881	1278989	1331115	821657	438259
1891	1187760	1172402	724774	383758
1901	1152829	1076188	646932	345874
1911	1162044	1035495	610984	331165
1926	1149092	969902	552907	300091
1936	1220411	942272	525468	280269
1946	1281117	917306	492797	263887
1951	1336576	898870	471895	253252
1956	1338942	877238	446221	235863
1961	1332149	849203	419465	217524
1966	1414415	859334	401950	208303
1971	1498140	882002	390902	207204
1979	1743861	979819	418500	226037
1981	1790521	998315	424410	230159
1986	1852649	1020577	431409	236008
1991	1860949	1009533	423031	232206
1996	1924702	1033903	433231	234251
2002	2105579	1100614	464296	246714
2006	2295123	1173340	504121	267264
2011	2504814	1246088	542547	294803



An Phríomh-Oifig Staidrimh
Central Statistics Office

Home Statistics **Databases** Methods About Us

You are here > Home / StatBank Ireland / This is Ireland Part 1

Theme

People and Society
Census of Population
2011 Census Results
This is Ireland Part 1
This is Ireland Part 2

Statistical Product - This is Ireland Part 1 **RSS**

Current Tables

CDD01 Population by Province or County, CensusYear and Sex (1841-2011) - Modified on 02/04/13 at 10:51
Download .px file (Size: 16.3 kb)
CDD02 Population by Sex, Age Group and CensusYear (1926-2011) - Modified on 29/03/12 at 9:59
Download .px file (Size: 2.9 kb)
CDD03 Population by Constituency, Statistic and CensusYear (2006-2011) - Modified on 29/03/12 at 9:59
Download .px file (Size: 5.3 kb)

<http://www.cso.ie/px/pxeirestat/Statire/SelectVarVal/Define.asp?maintable=CDD01&PLanguage=0>

Opening an Excel File

```
library(gdata)
library(reshape)
library(ggplot2)

c <- read.xls("03 Matrices and Data Frames/CensusData.xlsx")
```

```
> str(c)
'data.frame':  24 obs. of  5 variables:
 $ Year      : int  1841 1851 1861 1871 188
 $ Leinster: int  1973731 1672738 1457635
92 1220411 ...
 $ Munster  : int  2396161 1857736 1513558
2 942272 ...
 $ Connacht: int  1418859 1010031 913135
68 ...
 $ Ulster   : int  740048 571052 517783 47
```


	Year ↕	Leinster ↕	Munster ↕	Connacht ↕	Ulster ↕
1	1841	1973731	2396161	1418859	740048
2	1851	1672738	1857736	1010031	571052
3	1861	1457635	1513558	913135	517783
4	1871	1339451	1393485	846213	474038
5	1881	1278989	1331115	821657	438259
6	1891	1187760	1172402	724774	383758

Data format inflexible

- Plotting in R
- Aggregating the data
- Processing data

Year	Leinster	Munster	Connacht	Ulster
1841	1973731	2396161	1418859	740048
1851	1672738	1857736	1010031	571052
1861	1457635	1513558	913135	517783
1871	1339451	1393485	846213	474038
1881	1278989	1331115	821657	438259
1891	1187760	1172402	724774	383758

	Year	Province	Population
1	1841	Leinster	1973731
2	1851	Leinster	1672738
3	1861	Leinster	1457635
4	1871	Leinster	1339451
5	1881	Leinster	1278989
6	1891	Leinster	1187760



melt() function – reshape library

`melt(data, id.vars, measure.vars)`

- `data` Data set to melt
- `id.vars` Id variables. If blank, will use all non `measure.vars` variables. Can be integer (variable position) or string (variable name)
- `measure.vars` Measured variables. If blank, will use all non `id.vars` variables. Can be integer (variable position) or string (variable name)

```
new<-melt(c,id.vars="Year",
          measure.vars=c("Leinster",
                          "Munster",
                          "Connacht",|
                          "Ulster"))

names(new)<-c("Year","Province","Population")
```

```
> head(new)
  Year variable  value
1 1841 Leinster 1973731
2 1851 Leinster 1672738
3 1861 Leinster 1457635
4 1871 Leinster 1339451
5 1881 Leinster 1278989
6 1891 Leinster 1187760
```

split()

- **split** divides the data in the vector x into the groups defined by f. The replacement forms replace values corresponding to such a division. unsplit reverses the effect of split.

```
> s<-split(new,new$Year)
```

```
> length(s)
```

```
[1] 24
```

```
> str(s)
```

```
List of 24
```

```
$ 1841:'data.frame':  4 obs. of  3 variables:
```

```
..$ Year      : int [1:4] 1841 1841 1841 1841
```

```
..$ Province  : Factor w/ 4 levels "Leinster","Munster",...: 1 2 3 4
```

```
..$ Population: int [1:4] 1973731 2396161 1418859 740048
```

```
$ 1851:'data.frame':  4 obs. of  3 variables:
```

```
..$ Year      : int [1:4] 1851 1851 1851 1851
```

```
..$ Province  : Factor w/ 4 levels "Leinster","Munster",...: 1 2 3 4
```

```
..$ Population: int [1:4] 1672738 1857736 1010031 571052
```

aggregate()

```
a<-aggregate(Population~Year,new,sum)
```

```
> head(a)
```

	Year	Population
1	1841	6528799
2	1851	5111557
3	1861	4402111
4	1871	4053187
5	1881	3870020
6	1891	3468694

```
> tail(a)
```

	Year	Population
19	1986	3540643
20	1991	3525719
21	1996	3626087
22	2002	3917203
23	2006	4239848
24	2011	4588252

qplot()

```
qplot(data=new, x=Year, y=Population, color=Province, geom="line")
```

