

CT5102: Programming for Data Analytics

Week 8: Tidy Data and dplyr

<https://github.com/JimDuggan/CT5102>

Dr. Jim Duggan,
Information Technology,
School of Engineering & Informatics

Key Reference



Journal of Statistical Software

August 2014, Volume 59, Issue 10.

<http://www.jstatsoft.org/>

Tidy Data

Hadley Wickham
RStudio

Abstract

A huge amount of effort is spent cleaning data to get it ready for analysis, but there has been little research on how to make data cleaning as easy and effective as possible. This paper tackles a small, but important, component of data cleaning: data tidying. Tidy datasets are easy to manipulate, model and visualize, and have a specific structure: each variable is a column, each observation is a row, and each type of observational unit is a table. This framework makes it easy to tidy messy datasets because only a small set of tools are needed to deal with a wide range of un-tidy datasets. This structure also makes it easier to develop tidy tools for data analysis, tools that both input and output tidy datasets. The advantages of a consistent data structure and matching tools are demonstrated with a case study free from mundane data manipulation chores.

Keywords: data cleaning, data tidying, relational databases, R.

Overview

- What is data tidying?
 - Structuring datasets to facilitate analysis
- The tidy data standard is designed to:
 - Facilitate initial exploration and analysis of data
 - Simplify the development of data analysis tools that work well together
- Principles closely related to relational algebra (Codd 1990)
- Related packages: ggplot2, reshape, reshape2, plyr, dplyr

Typical Structure: Rows and Columns (Wickham 2014)

	treatmenta	treatmentb
John Smith	—	2
Jane Doe	16	11
Mary Johnson	3	1

Table 1: Typical presentation dataset.

	John Smith	Jane Doe	Mary Johnson
treatmenta	—	16	3
treatmentb	2	11	1

Table 2: The same data as in Table 1 but structured differently.

Numbers refer to the result of the treatments on a given person.

Data Semantics

- A dataset is a collection of values, usually numbers (if quantitative) or strings (if qualitative)
- Every value belongs to a variable and an observation
- An observation contains all values measured on the same unit (e.g. person or day)

	treatmenta	treatmentb
John Smith	—	2
Jane Doe	16	11
Mary Johnson	3	1

- The variables are:
 - Person (John, Jane, and Mary)
 - Treatment (a or b)
 - Result (6 values including NA)
- Every combination of person and treatment were measured.

Reorganising Tables

- Makes values, variables and observations clearer
- Variables in Columns
- Observations in Rows

	treatmenta	treatmentb
John Smith	—	2
Jane Doe	16	11
Mary Johnson	3	1

person	treatment	result
John Smith	a	—
Jane Doe	a	16
Mary Johnson	a	3
John Smith	b	2
Jane Doe	b	11
Mary Johnson	b	1

Tidy Data

- In tidy data
 - Each variable forms a column
 - Each observation forms a row
 - Each type of observational unit forms a table

person	treatment	result
John Smith	a	—
Jane Doe	a	16
Mary Johnson	a	3
John Smith	b	2
Jane Doe	b	11
Mary Johnson	b	1

Melting (Wickham 2014)

row	a	b	c
A	1	4	7
B	2	5	8
C	3	6	9

(a) Raw data

row	column	value
A	a	1
B	a	2
C	a	3
A	b	4
B	b	5
C	b	6
A	c	7
B	c	8
C	c	9

(b) Molten data

Table 5: A simple example of melting. (a) is melted with one colvar, row, yielding the molten dataset (b). The information in each table is exactly the same, just stored in a different way.

melt() function – reshape library

`melt(data, id.vars, measure.vars)`

- `data` Data set to melt
- `id.vars` Id variables. If blank, will use all non `measure.vars` variables. Can be integer (variable position) or string (variable name)
- `measure.vars` Measured variables. If blank, will use all non `id.vars` variables. Can be integer (variable position) or string (variable name)

Ordering Variables

- Fixed Variables
 - Describe the experimental design
 - Known in advance
- Measured Variables
 - What is measured in the study
- Fixed variables come first

religion	<\$10k	\$10-20k	\$20-30k	\$30-40k	\$40-50k	\$50-75k
Agnostic	27	34	60	81	76	137
Atheist	12	27	37	52	35	70
Buddhist	27	21	30	34	33	58
Catholic	418	617	732	670	638	1116
Don't know/refused	15	14	15	11	10	35
Evangelical Prot	575	869	1064	982	881	1486
Hindu	1	9	7	9	11	34
Historically Black Prot	228	244	236	238	197	223
Jehovah's Witness	20	27	24	24	21	30
Jewish	19	19	25	25	30	95

religion	income	freq
Agnostic	<\$10k	27
Agnostic	\$10-20k	34
Agnostic	\$20-30k	60
Agnostic	\$30-40k	81
Agnostic	\$40-50k	76
Agnostic	\$50-75k	137
Agnostic	\$75-100k	122
Agnostic	\$100-150k	109
Agnostic	>150k	84
Agnostic	Don't know/refused	96

Example

Gender	0 - 14 years	15 - 24 years	25 - 44 years	45 - 64 years	65 years and over
Male	501,189	290,898	717,055	520,243	243,314
Female	478,401	289,352	733,085	522,636	292,079

```
c.mf <- read.xls("08 Tidy Data/Census2011MF.xlsx")  
  
names(c.mf) <- c("Gender", "0-14", "15-24", "25-44", "45-64", "65+")  
cn <- melt(c.mf, id.vars = "Gender")  
names(cn) <- c("Gender", "Cohort", "Population")
```

Result

> c.mf

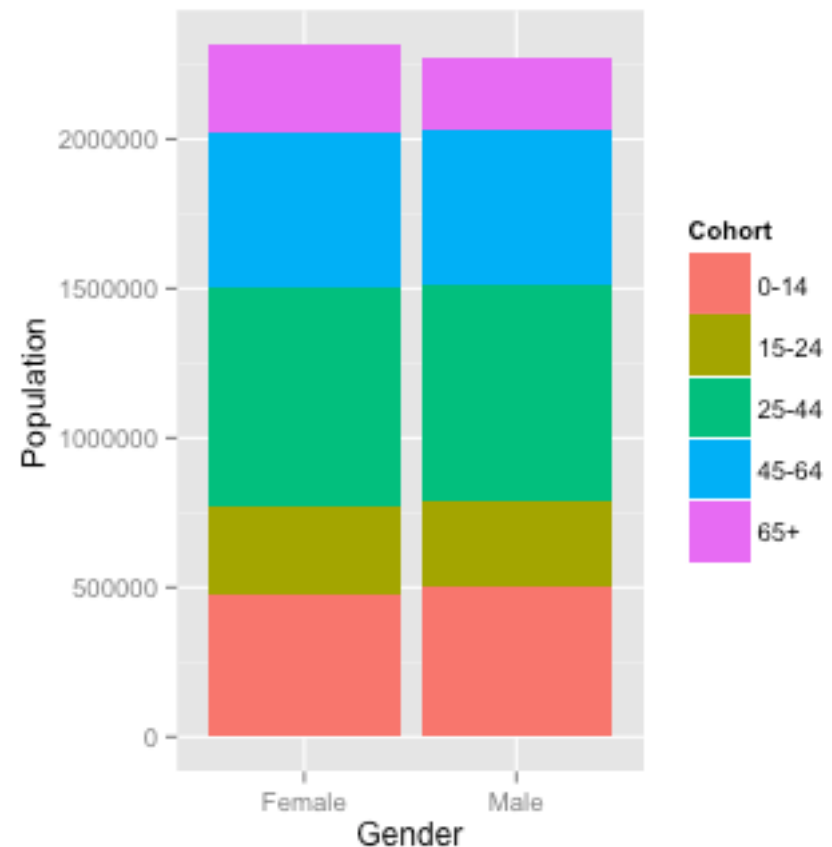
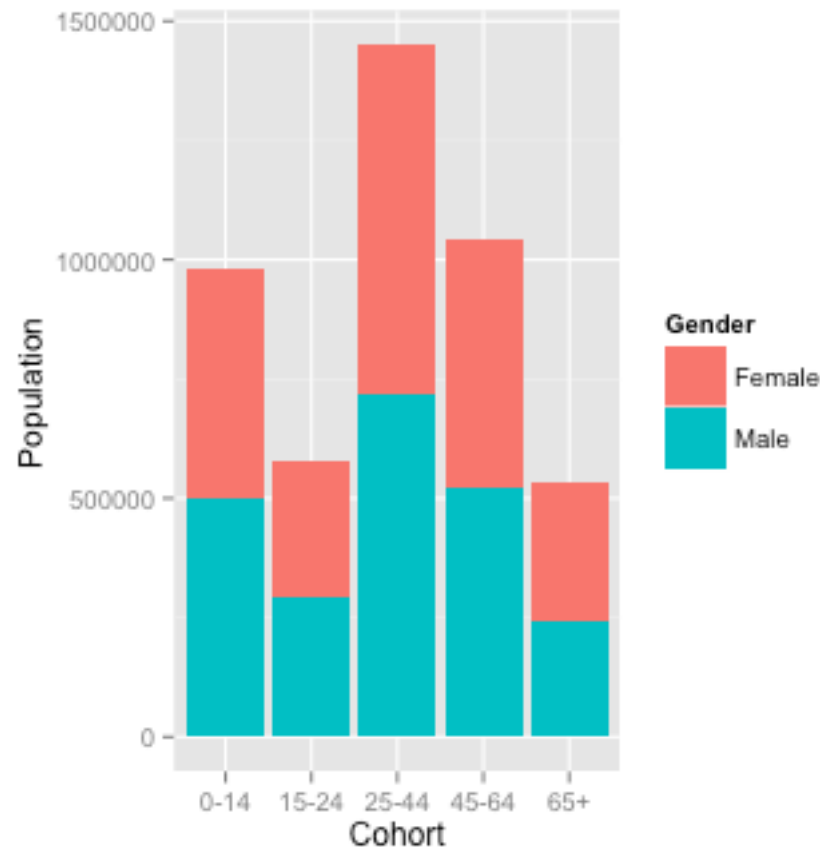
	Gender	0-14	15-24	25-44	45-64	65+
1	Male	501189	290898	717055	520243	243314
2	Female	478401	289352	733085	522636	292079

> cn

	Gender	Cohort	Population
1	Male	0-14	501189
2	Female	0-14	478401
3	Male	15-24	290898
4	Female	15-24	289352
5	Male	25-44	717055
6	Female	25-44	733085
7	Male	45-64	520243
8	Female	45-64	522636
9	Male	65+	243314
10	Female	65+	292079

Process Tidy Data with qplot

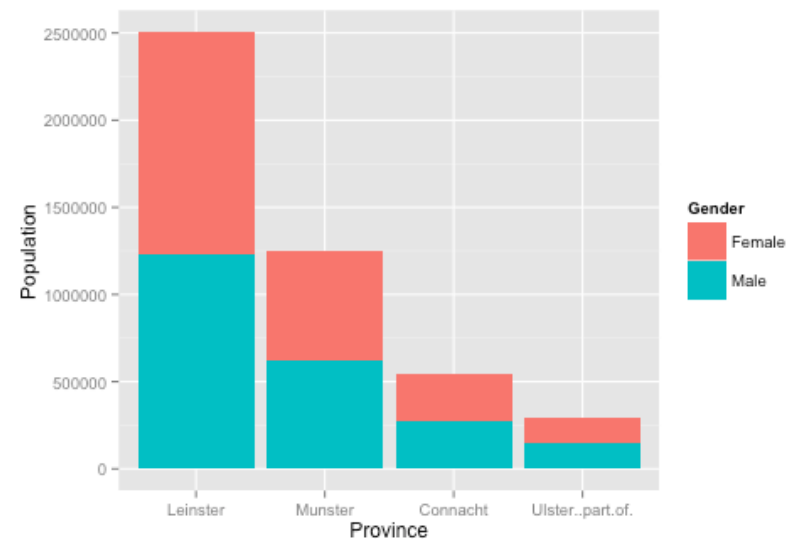
```
g1<-qplot(x=Cohort, y=Population, fill=Gender,data=cn, geom="bar", stat="identity")  
g2<-qplot(x=Gender, y=Population, fill=Cohort,data=cn, geom="bar", stat="identity")
```



Challenge 8.1

Gender	Leinster	Munster	Connacht	Ulster (part of)
Male	1,233,352	620,260	271,110	147,977
Female	1,271,462	625,828	271,437	146,826

- Show the above table in its tidy form
- Write the melt code to create a tidy data set
- Visualise the data using qplot



Multiple variables stored in one column (after melt)

Leinster_M	Munster_M	Connacht_M	Ulster_M	Leinster_F	Munster_F	Connacht_F	Ulster_F
1,233,352	620,260	271,110	147,977	1,271,462	625,828	271,437	146,826

```
c.col <- read.xls("08 Tidy Data/Census2011Combined.xlsx")  
c.c<-melt(c.col,id.vars = NULL)  
names(c.c)<-c("Province.Gender","Population")
```

```
> c.c  
  Province.Gender Population  
1   Leinster_M    1233352  
2   Munster_M     620260  
3   Connacht_M    271110  
4   Ulster_M      147977  
5   Leinster_F    1271462  
6   Munster_F     625828  
7   Connacht_F    271437  
8   Ulster_F      146826
```

Split a column into two

```
c.c$Province.Gender<-sub("_M","_Male",c.c$Province.Gender)  
c.c$Province.Gender<-sub("_F","_Female",c.c$Province.Gender)
```

```
r<-strsplit(as.character(c.c$Province.Gender),"_")
```

```
> str(r)
```

```
List of 8
```

```
$ : chr [1:2] "Leinster" "Male"  
$ : chr [1:2] "Munster" "Male"  
$ : chr [1:2] "Connacht" "Male"  
$ : chr [1:2] "Ulster" "Male"  
$ : chr [1:2] "Leinster" "Female"  
$ : chr [1:2] "Munster" "Female"  
$ : chr [1:2] "Connacht" "Female"  
$ : chr [1:2] "Ulster" "Female"
```


Split a column into two

```
mat <- matrix(unlist(r), ncol=2, byrow=TRUE)  
df  <- as.data.frame(mat)
```

```
k.c$Province<-df[,1]  
c.c$Gender<-df[,2]
```

> mat

	[,1]	[,2]
[1,]	"Leinster"	"Male"
[2,]	"Munster"	"Male"
[3,]	"Connacht"	"Male"
[4,]	"Ulster"	"Male"
[5,]	"Leinster"	"Female"
[6,]	"Munster"	"Female"
[7,]	"Connacht"	"Female"
[8,]	"Ulster"	"Female"

> df

	V1	V2
1	Leinster	Male
2	Munster	Male
3	Connacht	Male
4	Ulster	Male
5	Leinster	Female
6	Munster	Female
7	Connacht	Female
8	Ulster	Female

Result

```
> c.c
```

	Province.Gender	Population	Province	Gender
1	Leinster_Male	1233352	Leinster	Male
2	Munster_Male	620260	Munster	Male
3	Connacht_Male	271110	Connacht	Male
4	Ulster_Male	147977	Ulster	Male
5	Leinster_Female	1271462	Leinster	Female
6	Munster_Female	625828	Munster	Female
7	Connacht_Female	271437	Connacht	Female
8	Ulster_Female	146826	Ulster	Female

Remove old column

```
> c.c$Province.Gender<-NULL  
> c.c  
  Population Province Gender  
1    1233352  Leinster  Male  
2     620260  Munster  Male  
3     271110 Connacht  Male  
4     147977   Ulster  Male  
5    1271462  Leinster Female  
6     625828  Munster Female  
7     271437 Connacht Female  
8     146826   Ulster Female
```

ddply()

ddply {plyr}

R Documentation

Split data frame, apply function, and return results in a data frame.

Description

For each subset of a data frame, apply function then combine results into a data frame.

Usage

```
ddply(.data, .variables, .fun = NULL, ..., .progress = "none",  
      .inform = FALSE, .drop = TRUE, .parallel = FALSE, .paropts = NULL)
```

```
ddply(.data, .variables, .fun = NULL, ..., .progress = "none",  
      .inform = FALSE, .drop = TRUE, .parallel = FALSE, .paropts = NULL)
```

Arguments

- | | |
|-------------------------|---|
| <code>.fun</code> | function to apply to each piece |
| <code>...</code> | other arguments passed on to <code>.fun</code> |
| <code>.progress</code> | name of the progress bar to use, see create_progress_bar |
| <code>.parallel</code> | if TRUE, apply function in parallel, using parallel backend provided by foreach |
| <code>.paropts</code> | a list of additional options passed into the foreach function when parallel computation is enabled. This is important if (for example) your code relies on external data or packages: use the <code>.export</code> and <code>.packages</code> arguments to supply them so that all cluster nodes have the correct environment set up for computing. |
| <code>.inform</code> | produce informative error messages? This is turned off by default because it substantially slows processing speed, but is very useful for debugging |
| <code>.data</code> | data frame to be processed |
| <code>.variables</code> | variables to split data frame by, as as.quoted variables, a formula or character vector |
| <code>.drop</code> | should combinations of variables that do not appear in the input data be preserved (FALSE) or dropped (TRUE, default) |

Summarise function

Summarise a data frame.

Description

Summarise works in an analogous way to [mutate](#), except instead of adding columns to an existing data frame, it creates a new data frame. This is particularly useful in conjunction with [ddply](#) as it makes it easy to perform group-wise summaries.

Usage

```
summarise(.data, ...)
```

Arguments

- `.data` the data frame to be summarised
- `...` further arguments of the form `var = value`

Example

```
> cn
```

	Gender	Cohort	Population
1	Male	0-14	501189
2	Female	0-14	478401
3	Male	15-24	290898
4	Female	15-24	289352
5	Male	25-44	717055
6	Female	25-44	733085
7	Male	45-64	520243
8	Female	45-64	522636
9	Male	65+	243314
10	Female	65+	292079

```
> ddply(cn, .(Gender),  
+       summarize,  
+       Total= sum(Population))
```

	Gender	Total
1	Female	2315553
2	Male	2272699

```
> ddply(cn, .(Cohort),  
+       summarize,  
+       Total=sum(Population))
```

	Cohort	Total
1	0-14	979590
2	15-24	580250
3	25-44	1450140
4	45-64	1042879
5	65+	535393

Challenge 8.2


Use `ddply()` to get the total population by province.

```
> cp
```

	Gender	Province	Population
1	Male	Leinster	1233352
2	Female	Leinster	1271462
3	Male	Munster	620260
4	Female	Munster	625828
5	Male	Connacht	271110
6	Female	Connacht	271437
7	Male	Ulster..part.of.	147977
8	Female	Ulster..part.of.	146826

Analysing Tidy Data

- Tidy data is only worthwhile if it makes analysis easier
- Tidy tools
 - Take tidy data sets as inputs and return tidy data sets as outputs
- Tools cover:
 - Data manipulation
 - Visualisation
 - Modelling



dplyr package

- <https://rpubs.com/justmarkham/dplyr-tutorial>
- <https://www.youtube.com/watch?v=8SGif63VW6E>
- <https://www.youtube.com/watch?v=Ue08LVuk790>
- <http://renkun.me/pipeR-tutorial/Examples/dplyr.html>



Using dplyr

- Functions
 - **filter**: keep rows matching criteria
 - **select**: pick columns by name
 - **arrange**: reorder rows
 - **mutate**: add new variables
 - **summarise**: reduce variables to values
- Approach
 - First argument is a data frame
 - Subsequent arguments say what to do with data frame
 - Always return a data frame (Tidy Data approach)

1. filter()

Keep rows by matching criteria

```
> filter(cn, Gender=="Male")
```

	Gender	Cohort	Population
1	Male	0-14	501189
2	Male	15-24	290898
3	Male	25-44	717055
4	Male	45-64	520243
5	Male	65+	243314

```
> filter(cn, Gender=="Male" & Cohort=="0-14")
```

	Gender	Cohort	Population
1	Male	0-14	501189



```
> filter(cn, Gender=="Male", Cohort=="0-14")
```

	Gender	Cohort	Population
1	Male	0-14	501189

flights data

```
> str(flights)
Classes 'tbl_df', 'tbl' and 'data.frame':    227496 obs. of  14 variables:
 $ date      : Date, format: "2011-01-01" "2011-01-02" "2011-01-03" ...
 $ hour      : int   14 14 13 14 14 13 13 13 14 14 ...
 $ minute    : int    0 1 52 3 5 59 59 55 43 43 ...
 $ dep       : int  1400 1401 1352 1403 1405 1359 1359 1355 1443 1443 ...
 $ arr       : int  1500 1501 1502 1513 1507 1503 1509 1454 1554 1553 ...
 $ dep_delay: int    0 1 -8 3 5 -1 -1 -5 43 43 ...
 $ arr_delay: int   -10 -9 -8 3 -3 -7 -1 -16 44 43 ...
 $ carrier   : chr   "AA" "AA" "AA" "AA" ...
 $ flight    : int  428 428 428 428 428 428 428 428 428 428 ...
 $ dest      : chr   "DFW" "DFW" "DFW" "DFW" ...
 $ plane     : chr   "N576AA" "N557AA" "N541AA" "N403AA" ...
 $ cancelled: int    0 0 0 0 0 0 0 0 0 0 ...
 $ time      : int   40 45 48 39 44 45 43 40 41 45 ...
 $ dist      : int   224 224 224 224 224 224 224 224 224 224 ...
```

```
sfo <- filter(flights, dest == "SFO")  
qplot(date, dep_delay, data = sfo)  
qplot(date, arr_delay, data = sfo)  
qplot(arr_delay, dep_delay, data = sfo)
```



2. select() *pick columns by name*

```
> select(cn, -Cohort)
```

	Gender	Population
1	Male	501189
2	Female	478401
3	Male	290898
4	Female	289352
5	Male	717055
6	Female	733085
7	Male	520243
8	Female	522636
9	Male	243314
10	Female	292079

```
> select(cn, Gender, Population)
```

	Gender	Population
1	Male	501189
2	Female	478401
3	Male	290898
4	Female	289352
5	Male	717055
6	Female	733085
7	Male	520243
8	Female	522636
9	Male	243314
10	Female	292079

select()

pick columns by name

```
> select(cn, contains("Pop"))
```

	Population
1	501189
2	478401
3	290898
4	289352
5	717055
6	733085
7	520243
8	522636
9	243314
10	292079

```
> select(cn, starts_with("Coh"))
```

	Cohort
1	0-14
2	0-14
3	15-24
4	15-24
5	25-44
6	25-44
7	45-64
8	45-64
9	65+
10	65+

Special functions with select()

Special functions

As well as using existing functions like `:` and `c`, there are a number of special functions that only work inside `select`

- `starts_with(x, ignore.case = TRUE)`: names starts with `x`
- `ends_with(x, ignore.case = TRUE)`: names ends in `x`
- `contains(x, ignore.case = TRUE)`: selects all variables whose name contains `x`
- `matches(x, ignore.case = TRUE)`: selects all variables whose name matches the regular expression `x`
- `num_range("x", 1:5, width = 2)`: selects all variables (numerically) from `x01` to `x05`.
- `one_of("x", "y", "z")`: selects variables provided in a character vector.
- `everything()`: selects all variables.

3. arrange() *reorder rows*

```
> arrange(cn, desc(Population))
```

	Gender	Cohort	Population
1	Female	25-44	733085
2	Male	25-44	717055
3	Female	45-64	522636
4	Male	45-64	520243
5	Male	0-14	501189
6	Female	0-14	478401
7	Female	65+	292079
8	Male	15-24	290898
9	Female	15-24	289352
10	Male	65+	243314

```
> arrange(cn, Cohort, desc(Population))
```

	Gender	Cohort	Population
1	Male	0-14	501189
2	Female	0-14	478401
3	Male	15-24	290898
4	Female	15-24	289352
5	Female	25-44	733085
6	Male	25-44	717055
7	Female	45-64	522636
8	Male	45-64	520243
9	Female	65+	292079
10	Male	65+	243314

4. mutate()

Add new variables

```
> mutate(cn, Percentage=Population/sum(Population))
```

	Gender	Cohort	Population	Percentage
1	Male	0-14	501189	0.10923310
2	Female	0-14	478401	0.10426650
3	Male	15-24	290898	0.06340062
4	Female	15-24	289352	0.06306367
5	Male	25-44	717055	0.15628065
6	Female	25-44	733085	0.15977435
7	Male	45-64	520243	0.11338588
8	Female	45-64	522636	0.11390743
9	Male	65+	243314	0.05302978
10	Female	65+	292079	0.06365801

5. summarise()

Reduce variables to values

```
> summarise(cn, total=sum(Population))
  total
1 4588252
> by_gender<-group_by(cn, Gender)
> summarise(by_gender, total=sum(Population))
Source: local data frame [2 x 2]
```

	Gender	total
	(fctr)	(int)
1	Female	2315553
2	Male	2272699

.

Summary Functions

- `min(x)`, `median(x)`, `max(x)`,
- `quantile(x, p)`
- `n()`, `n_distinct()`, `sum(x)`, `mean(x)`
- `sum(x > 10)`, `mean(x > 10)`
- `sd(x)`, `var(x)`, `iqr(x)`, `mad(x)`

Pipe Operator in R %>%

- $x \%>\% f(y) \rightarrow f(x, y)$

```
> ans<-1:5 %>% sqrt() %>% + 20
```

```
> ans
```

```
[1] 21.00000 21.41421 21.73205 22.00000 22.23607
```

```
> filter(cn,Gender=="Male") %>% filter(Cohort=="0-14")
```

```
Gender Cohort Population Percentage
```

```
1 Male 0-14 501189 0.1092331
```

```
> filter(cn,Gender=="Male") %>% filter(Population==max(Population))
```

```
Gender Cohort Population Percentage
```

```
1 Male 25-44 717055 0.1562806
```

Challenge 8.3

Constituency	Population	Seats	Population Per Day Member	Constituency	Population	Seats	Population Per Day Member
Carlow-Kilkenny	145,659	5	29,132	Dún Laoghaire	105,029	4	26,257
Cavan-Monaghan	133,666	5	26,733	Galway East	110,085	4	27,521
Clare	111,336	4	27,834	Galway West	140,568	5	28,114
Cork East	114,365	4	28,591	Kerry North-West Limerick	80,883	3	26,961
Cork North-Central	104,911	4	26,228	Kerry South	77,971	3	25,990
Cork North-West	81,545	3	27,182	Kildare North	120,048	4	30,012
Cork South-Central	135,259	5	27,052	Kildare South	90,264	3	30,088
Cork South-West	82,952	3	27,651	Laois-Offaly	152,825	5	30,565
Donegal North-East	82,824	3	27,608	Limerick City	102,638	4	25,660
Donegal South-West	78,313	3	26,104	Limerick	81,679	3	27,226
Dublin Central	113,792	4	28,448	Longford-Westmeath	116,802	4	29,200
Dublin Mid-West	110,427	4	27,607	Louth	143,272	5	28,654
Dublin North	114,322	4	28,580	Mayo	130,638	5	26,128
Dublin North-Central	74,501	3	24,834	Meath East	86,572	3	28,857
Dublin North-East	81,560	3	27,187	Meath West	85,550	3	28,517
Dublin North-West	79,028	3	26,343	Roscommon - South Leitrim	80,973	3	26,991
Dublin South	140,543	5	28,109	Sligo-North Leitrim	80,283	3	26,761
Dublin South-Central	127,223	5	25,445	Tipperary North	85,024	3	28,341
Dublin South-East	103,833	4	25,958	Tipperary South	79,748	3	26,583
Dublin South-West	105,597	4	26,399	Waterford	112,198	4	28,050
Dublin West	117,214	4	29,304	Wexford	145,320	5	29,064
				Wicklow	141,012	5	28,202

Using tbl_df()

```
> con <- tbl_df(read.xls("08 Tidy Data/Constituencies.xlsx"))
```

```
> con
```

Source: local data frame [43 x 4]

	Constituency (fctr)	Population (int)	Seats (int)	Population.Per.Dail.Member (int)
1	Carlow-Kilkenny	145659	5	29132
2	Cavan-Monaghan	133666	5	26733
3	Clare	111336	4	27834
4	Cork East	114365	4	28591
5	Cork North-Central	104911	4	26228
6	Cork North-West	81545	3	27182
7	Cork South-Central	135259	5	27052
8	Cork South-West	82952	3	27651
9	Donegal North-East	82824	3	27608
10	Donegal South-West	78313	3	26104
..

Queries

1. List all the three seat constituencies
2. List all constituencies in Dublin
3. Show table without Population Per Dail Member
4. Add in the proportion of population for each constituency
5. Arrange data set by population and number of seats (desc)
6. Show total population covered by the different constituency types
7. Show the percentages for query (6).

dplyr assignment



An Phríomh-Oifig Staidrimh
Central Statistics Office

Gaeilge

 Search Statbank

Home

Statistics

Databases

Methods

About Us

You are here > [Home](#) / StatBank Ireland / House Prices

Theme

- People and Society
 - Census of Population
 - Annual Population Estimates

Statistical Product - House Prices [RSS](#)

Key Tables

- HPM01 Residential Property Price Index by Month and Type of Residential Property (2005M01-2015M08) - Modified on 30/09/15 at 12:10
[Download .px file](#) (Size: 10.4 kb)

Year.Month	National - all residential properties	Dublin - all residential properties	National excluding Dublin - all residential properties	National - houses	Dublin - houses	National excluding Dublin - houses	National - apartments	Dublin - apartments
2005M01	100.0	100.0	100.0	100.0	100.0	100.0	100.0	100.0
2005M02	100.4	101.1	100.1	100.5	100.8	100.3	100.1	101.7
2005M03	100.6	101.2	100.3	100.8	101.2	100.6	99.9	101.3
2005M04	101.3	102.2	100.9	101.6	102.4	101.2	99.7	101.5
2005M05	102.0	102.8	101.5	102.3	103.4	101.9	99.7	101.1
2005M06	102.9	103.5	102.6	103.4	104.2	103.1	99.6	101.5
2005M07	104.3	104.7	104.0	105.0	105.6	104.7	100.0	102.0