

Programming for Data Analytics

3. Matrices and Data Frames

Dr. Jim Duggan,
School of Engineering & Informatics
National University of Ireland Galway.

https://twitter.com/_jimduggan



Course Overview

Lectures 1-3	R Fundamentals <i>Atomic Vectors – Functions – Lists – Matrices – Data Frames</i>
Lectures 4-9	Data Science with R <i>ggplot2 – dplyr – tidyr – stringr – lubridate – purrr</i>
Lectures 10-11	Advanced Programming with R <i>Environments – Closures – S3 Object System</i>
Lectures 12	Machine Learning with R – Case Studies <i>Electricity Generation, Health</i>



Lecture Overview

- Matrices
 - Subsetting
 - apply function
- Data Frames
 - Filtering
 - Adding Columns
 - Cleaning Data

Lectures
1-3

R Fundamentals

*Atomic Vectors – Functions – Lists – **Matrices** – **Data Frames***

Lectures
4-9

Data Science with R

ggplot2 – dplyr – tidyr – stringr – lubridate – purrr

Lectures
10-11

Advanced Programming with R

Environments – Closures – S3 Object System

Lectures
12

Machine Learning with R – Case Studies

Electricity Generation, Health



(1) Matrices

	Homogenous	Heterogenous
1d	Atomic Vector	List
2d	Matrix	Data Frame
nd	Array	

- A matrix can be initialized from a vector, where the numbers of rows and columns are specified.
- R stores matrices by column-major order, and by default matrices are filled in this manner.

Declaring a matrix

```
>  
> a <- matrix(1:6, ncol=3, nrow=2)  
>  
> a  
      [,1] [,2] [,3]  
[1,]    1    3    5  
[2,]    2    4    6  
>  
> dim(a)  
[1] 2 3
```

Adding rows and columns

```
>  
> cbind(a,c(7,8))  
      [,1] [,2] [,3] [,4]  
[1,]    1    3    5    7  
[2,]    2    4    6    8
```

```
>  
>  
> rbind(a,c(7,8,9))  
      [,1] [,2] [,3]  
[1,]    1    3    5  
[2,]    2    4    6  
[3,]    7    8    9
```

Naming rows and columns

```
> rownames(a) <- c("A", "B")
>
> a
  [,1] [,2] [,3]
A     1     3     5
B     2     4     6
>
> colnames(a) <- c("a", "b", "c")
>
> a
  a b c
A 1 3 5
B 2 4 6
```

Subsetting Matrices

- The most common way of subsetting 2d matrix is a simple generalisation of 1d subsetting
- Supply a 1d index for each dimension, separated by a comma
- Blank subsetting is useful, as it lets you keep all rows or all columns

Using row index...

```
> b <- matrix(1:9, nrow=3)
>
> colnames(b) <- c("A", "B", "C")
>
> b
```

	A	B	C
[1,]	1	4	7
[2,]	2	5	8
[3,]	3	6	9

```
>
> b[1:2,]
```

	A	B	C
[1,]	1	4	7
[2,]	2	5	8

```
> b[c(T,F),]
```

	A	B	C
[1,]	1	4	7
[2,]	3	6	9

```
>
>
> b[-3,]
```

	A	B	C
[1,]	1	4	7
[2,]	2	5	8

Using column index...

```
> b
```

	A	B	C
[1,]	1	4	7
[2,]	2	5	8
[3,]	3	6	9

```
>
```

```
> b[,1:2]
```

	A	B
[1,]	1	4
[2,]	2	5
[3,]	3	6

```
> b[,c(T,F)]
```

	A	C
[1,]	1	7
[2,]	2	8
[3,]	3	9

```
>
```

```
> b[,c("A", "C")]
```

	A	C
[1,]	1	7
[2,]	2	8
[3,]	3	9

Sample Matrix Operations

Operator or Function	Description
$A * B$	Element-wise multiplication
A / B	Element-wise division
$A \%*\% B$	Matrix multiplication
$t(A)$	Transpose of A
$e \leftarrow \text{eigen}(A)$	List of eigenvalues and eigenvectors for matrix A



apply() function

- The **apply()** function can be used to process rows and columns for a matrix, and the general form of this function (Matloff 2009) is **apply(m, dimcode, f, fargs)**, where:
 - **m** is the target matrix
 - **dimcode** identifies whether it's a row or column target. The number 1 applies to rows, whereas 2 applies to columns
 - **f** is the function to be called
 - **fargs** are the optional set of arguments that can be applied to the function **f**.



Examples...

```
> b
```

	A	B	C
[1,]	1	4	7
[2,]	2	5	8
[3,]	3	6	9

```
>
```

```
> apply(b,1,sum)
```

[1]	12	15	18
-----	----	----	----

```
> b
```

	A	B	C
[1,]	1	4	7
[2,]	2	5	8
[3,]	3	6	9

```
>
```

```
> apply(b,2,sum)
```

A	B	C
6	15	24

(2) Data Frames

- The most common way of storing data in R
- Under the hood, a data frame is a list of equal-length vectors
- A two-dimensional structure, it shares properties of both a list and a matrix

	Homogenous	Heterogenous
1d	Atomic Vector	List
2d	Matrix	Data Frame
nd	Array	

Creating a data frame...

```
> df <- data.frame(x=1:5,y=LETTERS[1:5],stringsAsFactors=F)
```

```
>
```

```
> str(df)
```

```
'data.frame':  5 obs. of  2 variables:
```

```
$ x: int  1 2 3 4 5
```

```
$ y: chr  "A" "B" "C" "D" ...
```

```
> df
```

	x	y
1	1	A
2	2	B
3	3	C
4	4	D
5	5	E

Sample Data Set (mtcars)

A data frame with 32 observations on 11 variables.

- [, 1] mpg Miles/(US) gallon
- [, 2] cyl Number of cylinders
- [, 3] disp Displacement (cu.in.)
- [, 4] hp Gross horsepower
- [, 5] drat Rear axle ratio
- [, 6] wt Weight (1000 lbs)
- [, 7] qsec 1/4 mile time
- [, 8] vs V/S
- [, 9] am Transmission (0 = automatic, 1 = manual)
- [,10] gear Number of forward gears
- [,11] carb Number of carburetors



Using head()

```
> head(mtcars)
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
Mazda RX4	21.0	6	160	110	3.90	2.620	16.46	0	1	4	4
Mazda RX4 Wag	21.0	6	160	110	3.90	2.875	17.02	0	1	4	4
Datsun 710	22.8	4	108	93	3.85	2.320	18.61	1	1	4	1
Hornet 4 Drive	21.4	6	258	110	3.08	3.215	19.44	1	0	3	1
Hornet Sportabout	18.7	8	360	175	3.15	3.440	17.02	0	0	3	2
Valiant	18.1	6	225	105	2.76	3.460	20.22	1	0	3	1

Subsetting rows (note list operator \$ can be used to identify column)

```
> mtcars[mtcars$gear == 5,]
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
Porsche 914-2	26.0	4	120.3	91	4.43	2.140	16.7	0	1	5	2
Lotus Europa	30.4	4	95.1	113	3.77	1.513	16.9	1	1	5	2
Ford Pantera L	15.8	8	351.0	264	4.22	3.170	14.5	0	1	5	4
Ferrari Dino	19.7	6	145.0	175	3.62	2.770	15.5	0	1	5	6
Maserati Bora	15.0	8	301.0	335	3.54	3.570	14.6	0	1	5	8

```
>
```

```
> mtcars[mtcars$gear == 5 & mtcars$cyl == 8,]
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
Ford Pantera L	15.8	8	351	264	4.22	3.17	14.5	0	1	5	4
Maserati Bora	15.0	8	301	335	3.54	3.57	14.6	0	1	5	8

Accessing rows via indices

```
> mtcars[1:8,]
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
Mazda RX4	21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	4
Mazda RX4 Wag	21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	4
Datsun 710	22.8	4	108.0	93	3.85	2.320	18.61	1	1	4	1
Hornet 4 Drive	21.4	6	258.0	110	3.08	3.215	19.44	1	0	3	1
Hornet Sportabout	18.7	8	360.0	175	3.15	3.440	17.02	0	0	3	2
Valiant	18.1	6	225.0	105	2.76	3.460	20.22	1	0	3	1
Duster 360	14.3	8	360.0	245	3.21	3.570	15.84	0	0	3	4
Merc 240D	24.4	4	146.7	62	3.69	3.190	20.00	1	0	4	2

Sampling from a data frame...

- Selecting n random observations from a data frame

```
> mtcars[sample(nrow(mtcars),2),]
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
Chrysler Imperial	14.7	8	440.0	230	3.23	5.345	17.42	0	0	3	4
Merc 280	19.2	6	167.6	123	3.92	3.440	18.30	1	0	4	4

```
>
```

```
> mtcars[sample(nrow(mtcars),2),]
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
Merc 240D	24.4	4	146.7	62	3.69	3.190	20.0	1	0	4	2
Toyota Corolla	33.9	4	71.1	65	4.22	1.835	19.9	1	1	4	1



Filtering rows and columns

```
> mtcars[mtcars$cyl==6,c("mpg","cyl")]
```

	mpg	cyl
Mazda RX4	21.0	6
Mazda RX4 Wag	21.0	6
Hornet 4 Drive	21.4	6
Valiant	18.1	6
Merc 280	19.2	6
Merc 280C	17.8	6
Ferrari Dino	19.7	6

Filtering behaviour (> 1 column)

```
> mtcars[1:4,1:2]
```

	mpg	cyl
Mazda RX4	21.0	6
Mazda RX4 Wag	21.0	6
Datsun 710	22.8	4
Hornet 4 Drive	21.4	6

```
> str(mtcars[1:4,1:2])
```

```
'data.frame':  4 obs. of  2 variables:  
 $ mpg: num  21 21 22.8 21.4  
 $ cyl: num  6 6 4 6
```

Filtering behaviour (1 column)

- This shows the difference between **simplifying** and **preserving**
- This also is a common source of programming errors, as R automatically converts a 1 column data frame result to a vector
- Use `drop=FALSE` to prevent
- tibbles (later) are designed to avoid this problem.

```
> mtcars[1:4,1]
[1] 21.0 21.0 22.8 21.4
```

```
>
```

```
> mtcars[1:4,1,drop=F]
      mpg
Mazda RX4      21.0
Mazda RX4 Wag  21.0
Datsun 710     22.8
Hornet 4 Drive 21.4
```



Challenge 3.1

	mpg	cyl	displacement	hp	drat	wt	qsec	vs	am	gear	carb
Mazda RX4	21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	4
Mazda RX4 Wag	21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	4
Datsun 710	22.8	4	108.0	93	3.85	2.320	18.61	1	1	4	1
Hornet 4 Drive	21.4	6	258.0	110	3.08	3.215	19.44	1	0	3	1
Hornet Sportabout	18.7	8	360.0	175	3.15	3.440	17.02	0	0	3	2
Valiant	18.1	6	225.0	105	2.76	3.460	20.22	1	0	3	1
Duster 360	14.3	8	360.0	245	3.21	3.570	15.84	0	0	3	4
Merc 240D	24.4	4	146.7	62	3.69	3.190	20.00	1	0	4	2
Merc 230	22.8	4	140.8	95	3.92	3.150	22.90	1	0	4	2
Merc 280	19.2	6	167.6	123	3.92	3.440	18.30	1	0	4	4

- List all the cars that have an mpg greater than the average
- List the car(s) with the greatest displacement

Sample Data Set (mtcars)


A data frame with 32 observations on 11 variables.

- [, 1] mpg Miles/(US) gallon
- [, 2] cyl Number of cylinders
- [, 3] disp Displacement (cu.in.)
- [, 4] hp Gross horsepower
- [, 5] drat Rear axle ratio
- [, 6] wt Weight (1000 lbs)
- [, 7] qsec 1/4 mile time
- [, 8] vs V/S
- [, 9] am Transmission (0 = automatic, 1 = manual)
- [,10] gear Number of forward gears
- [,11] carb Number of carburetors



Adding Columns

- Often the initial data set may not contain sufficient information for analysis
- Adding new variables (columns) is an important feature to have
- Data frames support this: columns can be combined or new information used



Adding new columns

- Create a new column that contains kilometers per gallon

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
Mazda RX4	21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	4
Mazda RX4 Wag	21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	4
Datsun 710	22.8	4	108.0	93	3.85	2.320	18.61	1	1	4	1
Hornet 4 Drive	21.4	6	258.0	110	3.08	3.215	19.44	1	0	3	1
Hornet Sportabout	18.7	8	360.0	175	3.15	3.440	17.02	0	0	3	2
Valiant	18.1	6	225.0	105	2.76	3.460	20.22	1	0	3	1
Duster 360	14.3	8	360.0	245	3.21	3.570	15.84	0	0	3	4
Merc 240D	24.4	4	146.7	62	3.69	3.190	20.00	1	0	4	2
Merc 230	22.8	4	140.8	95	3.92	3.150	22.90	1	0	4	2
Merc 280	19.2	6	167.6	123	3.92	3.440	18.30	1	0	4	4

Solution

```
> mtcars$kpg <- mtcars$mpg * 8/5
```

```
>
```

```
> mtcars[1:8,]
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb	kpg
Mazda RX4	21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	4	33.60
Mazda RX4 Wag	21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	4	33.60
Datsun 710	22.8	4	108.0	93	3.85	2.320	18.61	1	1	4	1	36.48
Hornet 4 Drive	21.4	6	258.0	110	3.08	3.215	19.44	1	0	3	1	34.24
Hornet Sportabout	18.7	8	360.0	175	3.15	3.440	17.02	0	0	3	2	29.92
Valiant	18.1	6	225.0	105	2.76	3.460	20.22	1	0	3	1	28.96
Duster 360	14.3	8	360.0	245	3.21	3.570	15.84	0	0	3	4	22.88
Merc 240D	24.4	4	146.7	62	3.69	3.190	20.00	1	0	4	2	39.04

Challenge 3.2

- Create a new column called “Model”, which contains the type of car. How is the model information currently stored?

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
Mazda RX4	21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	4
Mazda RX4 Wag	21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	4
Datsun 710	22.8	4	108.0	93	3.85	2.320	18.61	1	1	4	1
Hornet 4 Drive	21.4	6	258.0	110	3.08	3.215	19.44	1	0	3	1
Hornet Sportabout	18.7	8	360.0	175	3.15	3.440	17.02	0	0	3	2
Valiant	18.1	6	225.0	105	2.76	3.460	20.22	1	0	3	1
Duster 360	14.3	8	360.0	245	3.21	3.570	15.84	0	0	3	4
Merc 240D	24.4	4	146.7	62	3.69	3.190	20.00	1	0	4	2
Merc 230	22.8	4	140.8	95	3.92	3.150	22.90	1	0	4	2
Merc 280	19.2	6	167.6	123	3.92	3.440	18.30	1	0	4	4

What does this code do?

```
> my_mtcars <- mtcars
>
> set.seed(0)
>
> rows <- sample(nrow(my_mtcars),3)
> col  <- sample(ncol(my_mtcars),1)
>
> my_mtcars[rows,col] <- NA
>
> rows
[1] 29  9 12
>
> col
[1] 7
```

Resulting data frame (3 rows)

```
> my_mtcars[rows,]
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
Ford Pantera L	15.8	8	351.0	264	4.22	3.17	NA	0	1	5	4
Merc 230	22.8	4	140.8	95	3.92	3.15	NA	1	0	4	2
Merc 450SE	16.4	8	275.8	180	3.07	4.07	NA	0	0	3	3

Find rows that have incomplete cases in the data set

```
> !complete.cases(my_mtcars)
```

```
[1] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE TRUE FALSE FALSE  
[12] TRUE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE  
[23] FALSE FALSE FALSE FALSE FALSE FALSE TRUE FALSE FALSE FALSE
```

```
>
```

```
> my_mtcars[!complete.cases(my_mtcars),]
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
Merc 230	22.8	4	140.8	95	3.92	3.15	NA	1	0	4	2
Merc 450SE	16.4	8	275.8	180	3.07	4.07	NA	0	0	3	3
Ford Pantera L	15.8	8	351.0	264	4.22	3.17	NA	0	1	5	4

```
>
```

```
> rows
```

```
[1] 29 9 12
```


Approximating missing values.

```
> mean(my_mtcars$qsec)
[1] NA
>
> mean(my_mtcars$qsec, na.rm = T)
[1] 17.80552
```

- Replace any missing values in the qsec column with the overall average value

Changing the values...

```
> my_mtcars[rows,"qsec"] <- mean(my_mtcars$qsec,na.rm = T)
>
> my_mtcars[rows,]
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
Ford Pantera L	15.8	8	351.0	264	4.22	3.17	17.80552	0	1	5	4
Merc 230	22.8	4	140.8	95	3.92	3.15	17.80552	1	0	4	2
Merc 450SE	16.4	8	275.8	180	3.07	4.07	17.80552	0	0	3	3

Setting invalid values to NA

```
> my_mtcars <- mtcars
>
> rows <- sample(nrow(my_mtcars),3)
> col <- sample(ncol(my_mtcars),1)
>
> my_mtcars[rows,col] <- -9000
>
> col <- sample(ncol(my_mtcars),1)
>
> my_mtcars[rows,col] <- -1000
> my_mtcars[rows,]
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
Ford Pantera L	15.8	8	351.0	264	4.22	3.17	-9000	0	1	-1000	4
Merc 230	22.8	4	140.8	95	3.92	3.15	-9000	1	0	-1000	2
Merc 450SE	16.4	8	275.8	180	3.07	4.07	-9000	0	0	-1000	3

Removing invalid values...

```
> my_mtcars[rows,]
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
Ford Pantera L	15.8	8	351.0	264	4.22	3.17	-9000	0	1	-1000	4
Merc 230	22.8	4	140.8	95	3.92	3.15	-9000	1	0	-1000	2
Merc 450SE	16.4	8	275.8	180	3.07	4.07	-9000	0	0	-1000	3

```
>
```

```
> clean <- data.frame(apply(my_mtcars,2,  
+                           function(x)ifelse(x<0,NA,x)))
```

```
>
```

```
> clean[rows,]
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
Ford Pantera L	15.8	8	351.0	264	4.22	3.17	NA	0	1	NA	4
Merc 230	22.8	4	140.8	95	3.92	3.15	NA	1	0	NA	2
Merc 450SE	16.4	8	275.8	180	3.07	4.07	NA	0	0	NA	3



Challenge 3.3

- Extend the previous example so that the apply function returns the average of all valid values instead of NA.
- *Hint: Apply functions can have more than one statement, once the {} brackets are used.*

```
> my_mtcars[rows,]
```

		mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
Ford	Pantera L	15.8	8	351.0	264	4.22	3.17	-9000	0	1	-1000	4
Merc	230	22.8	4	140.8	95	3.92	3.15	-9000	1	0	-1000	2
Merc	450SE	16.4	8	275.8	180	3.07	4.07	-9000	0	0	-1000	3